TASK 3

What new improvements were introduced in Gemini 3.0?

- Maximizes reasoning depth. The model may take significantly longer to reach a first token, but the output will be more carefully reasoned.
- Gemini 3 introduces granular control over multimodal vision processing via the media_resolution parameter. Higher resolutions improve the model's ability to read fine text or identify small details, but increase token usage and latency
- Gemini 3 Pro Image lets you generate and edit images from text prompts. It uses reasoning to "think" through a prompt and can retrieve real-time data—such as weather forecasts or stock charts

How does Gemini 3.0 improve coding & automation workflows?

- Generates richer, production-ready code with improved frontend quality and UI components; excels at "zero-shot" tasks like simulating OS interfaces from a single prompt or creating interactive 3D games/visualizations.
- Strict multimodal function calling with thought signatures ensures seamless multi-turn reasoning; supports parallel calls (e.g., querying weather in multiple cities) and integrations like Google Search for real-time data in workflows
- Handles end-to-end workflows autonomously (e.g., inbox organization via Gmail integration or booking services), with granular controls like thinking_level: "high" for complex reasoning and default temperature 1.0 to prevent looping.

How does Gemini 3.0 improve multimodal understanding?

- Supports up to 1M input tokens , allowing processing of entire videos, large PDFs, or full code repositories alongside text.
- New media_resolution parameter (low/medium/high) optimizes vision tasks—high for detailed OCR (1120 tokens/image), medium for PDFs (560 tokens), low for speed.

Name any two developer tools introduced with Gemini 3.0.

- Google Antigravity: A revolutionary agentic development platform that lets AI agents autonomously plan, code, test, and deploy full applications from natural language prompts